

Shaurya Mittal

+1 (647) 961-7749 | shaurya.mittal@uwaterloo.ca | shauryamittal.net | [GitHub](https://github.com/shaurmittal) | [LinkedIn](https://www.linkedin.com/in/shauryamittal)

EDUCATION

University of Waterloo

Bachelor of Computer Science (Honours) Co-op Program

2025 – Present

Waterloo, ON

SKILLS

Languages: Python, TypeScript, JavaScript, Go, Java, C, SQL

AI & LLMs: Claude API, OpenAI API, RAG, LLM Evaluation (LLM-as-judge), Vector Databases (Qdrant), MCP

Frameworks & Libraries: React, Next.js, Node.js, Express, Streamlit

Cloud & Data: AWS (S3, Athena, CloudWatch), Snowflake, PostgreSQL, MongoDB, Firestore, Firebase

Tools: Git, Docker, GitLab CI/CD, REST APIs, Stripe API, Chrome Extensions (Manifest V3)

EXPERIENCE

AI Solutions Engineering Intern

Docebo Inc.

May 2026 – August 2026

Toronto, ON

- Built a full-stack agent feedback dashboard (React, Glean OAuth, AWS CloudWatch/S3/Athena) giving agent owners role-scoped, self-service access to user feedback, replacing manual data pulls by AI Operations.
- Engineered a company-wide AI portal (Next.js, MCP server, GitLab-backed) serving 900+ employees as a single hub for AI guides, the full agent catalog with Glean source links, and self-service skill and agent registration.
- Diagnosed a silent pagination bug truncating Glean API data, then shipped a reconciliation/soft-delete sync framework (capped at 30% deletions per run) that keeps an 800+ agent catalog in sync automatically.
- Audited the Glean agent catalog end-to-end, removing 670+ stale, orphaned, or duplicate agents (a 67% reduction), and authored the lifecycle/governance strategy adopted for ongoing catalog hygiene.

PROJECTS

hookfuzz – Chaos Testing CLI for Payment Webhooks

Go, Node.js, Express, Stripe API | github.com/shaurmittal/hookfuzz

September 2026

- Built a Go CLI that finds bugs in Stripe webhook handlers by delivering signed events twice, out of order, late, or after failures, then checking user-defined correctness rules against the app's final state.
- Implemented delta-debugging test shrinking with seeded, reproducible runs, reducing a 44-event failing delivery sequence to a minimal 2-event reproduction that replays exactly from a single seed.
- Validated the tool against a sample Express shop with 3 planted bugs (duplicate fulfillment, dropped refund, stale subscription overwrite): it caught all 3, with zero false positives over 200 seeds on the fixed version.
- Designed fail-safe checks so that misconfigurations such as a wrong signing secret or a mistyped rule stop with an error instead of reporting a false pass.

Verte – Sustainable Shopping Chrome Extension

TypeScript, React, Node.js, OpenAI API, Snowflake | devpost.com/software/verte-niek4b

September 2026

- Won 2nd place in the Pivot Track at PivotHacks 2026 with a team of 4, building a Manifest V3 Chrome extension that injects secondhand and lower-carbon alternatives inline on Amazon.ca and Best Buy product pages.
- Built a Node.js proxy using the OpenAI API to classify products and recommend alternatives, verifying every suggestion against live Amazon.ca listings to filter out hallucinated or discontinued items.
- Integrated a Snowflake knowledge base of cited embodied-carbon estimates and added voice search via OpenAI transcription with silence detection, persisting budget and location context across queries.

FastAPI RAG Q&A Engine

Python, Claude API, Qdrant, Streamlit, Docker | github.com/shaurmittal/rag-fastapi-eval

September 2026

- Built a cited-answer RAG system over FastAPI's docs and source code using Qdrant, local bge-small embeddings, a cross-encoder reranker, and Claude with native citations, deployed as a Streamlit demo.
- Ran a 4-configuration ablation (chunking strategy × reranking) on 51 questions mined from 613 closed GitHub issues; paired bootstrap CIs showed reranking cost $\sim 15\times$ latency for no significant gain.
- Found that a pooled, TREC-style relevance review recovered 28 correct answers the first-pass labels missed, raising hit@5 by ~ 0.10 , a larger effect than any pipeline change.
- Built an LLM-as-judge pipeline scoring claim-level faithfulness, correctness, and citation precision, validated against an independent rater ($\kappa = 0.89$), grading 204 answers for \$9 via the Message Batches API.